


Generalized Linear Models

Tim Book

Summit Consulting, LLC



`Summit.describe()`



- **Summit is a statistical and econometric consulting firm here in DC.**
- **Work with both government agencies and private companies.**
- **I work in our Litigation directorate. I use data science to help lawyers win cases.**
- **Summit is currently hiring. (www.summitllc.us)**



Table of Contents


- **“Recap” of OLS**
- **“Recap” of Logistic Regression**
- **Zoom out, take a look at GLMs in general**
 - Discuss model evaluation, inference
- **Discuss some common GLMs and their applications**
- **An interlude on model inference**
- **GLMs for Classification!**
- **Code-a-long**
- **Pot Pourri (if time)**



“Recap” of OLS

OLS At-A-Glance

- Ordinary Least Squares is the classic “line of best fit.”
- Given:
 - x-variable(s).
 - y-variable on a *continuous scale*.
- We find:
 - The slope estimate $\hat{\beta}$. (How do we interpret this again?)
 - The variance estimate σ^2 .
- But what do we mean by *best* fit?



- The traditional OLS has the following form:

$$y_i = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_p x_{pi} + \varepsilon_i$$

For each $i = 1, \dots, n$

- We also assume $\varepsilon_i \sim \text{iid } N(0, \sigma^2)$.
- So now we can think of each $y_i \sim \text{ind } N(\mathbf{x}_i^T \boldsymbol{\beta}, \sigma^2)$.
- Notice in this notation I eliminated the need for thinking about residuals.



Finding $\hat{\boldsymbol{\beta}}$

- Now, if each $y_i \sim N(\mathbf{x}_i^T \boldsymbol{\beta}, \sigma^2)$, we can use *maximum likelihood estimation* to find the “best” $\hat{\boldsymbol{\beta}}$:

$$\hat{\boldsymbol{\beta}} = \underset{\boldsymbol{\beta}}{\operatorname{argmax}} f(\mathbf{y}|\boldsymbol{\beta}) = \underset{\boldsymbol{\beta}}{\operatorname{argmax}} \log f(\mathbf{y}|\boldsymbol{\beta})$$

$$= \underset{\boldsymbol{\beta}}{\operatorname{argmax}} \left(c - \sum_{i=1}^n \frac{1}{2\sigma} (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2 \right)$$

$$= \underset{\boldsymbol{\beta}}{\operatorname{argmin}} \sum_{i=1}^n (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2 \quad \leftarrow \text{Now THAT looks familiar!}$$



So...

- So the “least squares” technique is actually just a maximum likelihood technique!
- The maximum likelihood technique relied heavily on our *normality assumption*.
- But what if our response variable y isn't normal? What if it's not even *continuous*?!



“Recap” Logistic Regression


Binary Response

- Let's first consider the times where our response variable y is binary (coded as 0 or 1).
- This can be any type of variable where there are *exactly two* possible outcomes: yes/no, male/female, republican/democrat, on/off, etc.



Binary Response

- What would happen if you tried to fit an OLS model to a binary response?



What's the problem?

Besides the mechanical issues, using OLS to model binary response suffers two major flaws:

1. The range of responses is *unbounded*.
2. A binary random variable is not *normally distributed*.



Solution

Given data \mathbf{x} (I leave it out of what follows for brevity), we can assume our binary response follows a Bernoulli distribution with success probability p :

$$y \sim \text{ind Ber}(p)$$

Notice that we can't actually *observe* p ! Can only observe y .

Nonetheless, let's try to get this to look like a linear model anyway. We need a function that maps p (which can only be in $(0,1)$) to the real line \mathbb{R} .



What values can this have?

p



What values can this have?

$$\frac{p}{1 - p}$$

And does this value have a name?



What values can this have?

$$\log \frac{p}{1-p}$$

This is also sometimes written $\text{logit } p$



Finally, a linear model!

$$y_i = \log \frac{p_i}{1 - p_i} = \mathbf{x}_i^T \boldsymbol{\beta}$$

Notes:

- Remember that p , like y , depends on \mathbf{x}
- I didn't add the $+\varepsilon_i$. Why not?



Rearranging...

Since

$$y_i = \text{logit } p_i = \log \frac{p_i}{1 - p_i} \in \mathbb{R}$$

We can rewrite as


$$p_i = \frac{e^{y_i}}{1 + e^{y_i}} = \frac{e^{\mathbf{x}_i^T \boldsymbol{\beta}}}{1 + e^{\mathbf{x}_i^T \boldsymbol{\beta}}} \in (0,1)$$

(This is sometimes called the “inverse logit” function.)



Does it work?

Thanks to Bernoulli
assumption, y-axis can now
be interpreted as the “ p -axis”!



Bounded!

Solving...

So, uh, how do we now find the β s? Since we know y is Bernoulli, let's try that *MLE* technique again...

$$\hat{\beta} = \operatorname{argmax}_{\beta} f(\mathbf{y}|\beta) = \operatorname{argmax}_{\beta} \log f(\mathbf{y}|\beta)$$

$$= \operatorname{argmax}_{\beta} \log \left(\prod p^{y_i} (1-p)^{1-y_i} \right)$$

= ... substituting and rearranging ...

$$= \operatorname{argmax}_{\beta} \sum y_i (\mathbf{x}_i^T \beta) - \sum \log(1 - e^{\mathbf{x}_i^T \beta}) \quad \leftarrow \text{Gross and impossible!}$$




For logistic regression (and all (most) other (basic) GLMs), the coefficients β are determined via maximum likelihood and *Iteratively Reweighted Least Squares (IRLS)*.

Hand-wavy reminder of what IRLS is:

IRLS = Newton's Method + Chain Rule

Interpretation

- **OLS:** *“For a one unit increase in x_j , we would expect a $\hat{\beta}_j$ increase in y , holding all other variables constant.”*
- **Logistic Reg:** 

Reminder: $\text{logit}(p) = \mathbf{x}^T \hat{\beta}$

Reminder: $\frac{p}{1-p} = e^{\mathbf{x}^T \hat{\beta}}$

Interpretation

- **OLS:** “For a one unit increase in x_j , we would expect a $\hat{\beta}_j$ increase in y , holding all other variables constant.”
- **Logistic Reg:** “For a one unit increase in x_j , we would expect the odds of [success] to **increase by a factor** of $e^{\hat{\beta}_j}$ ”
- **Why?**

$$\frac{p}{1-p} = e^{\hat{\beta}_0 + \hat{\beta}_1(x+1)} = e^{\hat{\beta}_0 + \hat{\beta}_1 x} e^{\hat{\beta}_1}$$



GLMs In General

GLMs have 3 components

- **Linear (Systematic) component:** The x-variables and corresponding slopes.
- **Random component:** The distributional assumption you make about where the “randomness” in your model comes from. (eg, the Bernoulli distribution in logistic regression)
- **Link component:** The mapping function that “links” the random component to the linear component. (eg, logit)



Linear (Systematic) Component

- This is the part you already know about. I won't spend a lot of time here.

$$stuff = \beta_0 + \underbrace{\beta_1 x_1 + \cdots + \beta_p x_p}_{\text{This}}$$



- **This is where you make your distributional assumption about your response.**
 - Earlier, we assumed our 0/1 response was *Bernoulli*. In general, we could have assumed our response was *binomial*, if that made sense to us.
 - In OLS, we assumed our response was *normal*. (Yes, OLS is really just a GLM!)



What else is out there?

- **Actually, any distribution that is a member of the *Aitken Exponential Family* can be used in a GLM!**

$$c(y, \phi) \exp \left\{ \frac{b(\theta)T(y) - A(\theta)}{a(\phi)} \right\}$$

Look ugly? It is. Here are some Cliffsnote:



Commonly used Aitken EFs:

- **Normal**
- **Bernoulli/Binomial**
- **Poisson**
- **Exponential/Gamma**
- **Multinomial (much more on this later!)**

Less common ones:

- **Inverse Normal**
- **Negative Binomial**



Link Component

- The link component links the random component to the linear component. It must be invertible and (usually) fit this criterion:

$$g: \mathcal{Y} \rightarrow \mathbb{R}$$

*Response space (eg,
(0,1) for logistic reg.)*

The range of $\mathbf{x}^T \boldsymbol{\beta}$



Example: Logistic Regression Link

- The most common link function for logistic regression is the *logit function*, which we saw earlier:

$$\text{logit}(p) = \log \frac{p}{1-p}$$

- Another common one is the *probit function*:


$$\text{probit}(p) = \Phi^{-1}(p)$$

- NOTE: If you use the probit function, it is *NOT* called “probistic regression” (citation: Ellie Kazar)



They both look the same, so why choose?

- Visually, the logit and probit links are very hard to tell apart. Choosing one versus the other probably won't affect your predictions very much.
- The difference, then, comes in *interpretation*, if you care about that.
- The logit allows you to discuss your probabilities in terms of *odds*.
- The probit allows you to discuss your probabilities in terms of *percentiles*.



Formalizing the Logistic Regression

In purest mathematical terms:

$$Y_i \sim \text{ind Bin}(1, p_i)$$

$$\text{logit}(p_i) = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_p x_{pi}$$



Other Common GLMs

Poisson Regression

- *Poisson regression* is very popular for modeling *count data*, especially when the counts are small.
- **For example: Modeling household size versus number of bedrooms, number bathrooms, etc**
 - Household sizes start at 1 person, most being around 4 or 5, but can go higher.
- **Poisson regression typically uses the *log link*.**
 - Why? Counts are typically very right-skew. The natural log can pull unlikely large numbers lower, making the overall distribution appear more linear.
- **You can also use the *square root link*, but it is less common.**
 - It is, however, the variance stabilizing link!



Formally, Poisson Regression:

$$Y_i \sim \text{ind Poi}(\mu_i)$$

$$\log(\mu_i) = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_p x_{pi}$$


Discuss: How is this different from just taking the log of the response in a usual regression?



Gamma Regression

- Although it is rarer than the others, *gamma regression* can be used when your response is *waiting times* between events in a *Poisson process*.

- Think: Length of time for a lightbulb to go out



- There are two competing links for gamma regression: the *log link* and the *inverse link*. How do you decide?
 - If the plot of y vs $\log(x)$ looks linear, use the log link
 - If the plot of y vs $1/x$ looks linear, use the inverse link.

Formally, Gamma Regression:

$$Y_i \sim \text{ind Gamma}(\mu_i, \nu)$$

$$1/\mu_i = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$$

Where $\phi = 1/\nu$ is the *dispersion parameter*, which is used for model inference (more on this in a minute).



The Boring Stuff (Model Inference)

- Just like in OLS, the slope parameters $\hat{\beta}$ obey the Central Limit Theorem.

- So, as $n \rightarrow \infty$,
$$\hat{\beta}_j \rightarrow N\left(\hat{\beta}_j, SE(\hat{\beta}_j)\right)$$



Y tho? We often want to remove variables that aren't "important" to prevent overfitting and multicollinearity. We then test:

$$H_0: \beta_j = 0$$

$$H_A: \beta_j \neq 0$$

using Z inference. Python output takes all of the math-work out of this for you.

Deviance is a big deal. Let's set up some notation:

$\tilde{\beta}$: The MLE for model parameters in the *fully saturated model* (ie, the perfectly fit model.)

$\hat{\beta}$: The MLE for model parameters in *your* model.

$l(\hat{\beta}), l(\tilde{\beta})$: The likelihood functions of these models



Then, the *deviance* is defined as follows:

$$D = 2\hat{\phi}[l(\tilde{\beta}) - l(\hat{\beta})] > 0$$

*By definition, has
highest likelihood*

*Must be lower, since you are
using less information*

So, the deviance sort of measures how much “information” you lose by fitting something besides the perfect model. That is, *lower deviance is better*. You can think of it as an MSE analog.

Recall $\hat{\phi}$ is from the Aitken EF form. Usually, you don’t need to worry about it:

	Normal	Poisson	Bernoulli	Gamma
ϕ	σ^2	1	1	$1/\nu$

Deviance

However, the deviance doesn't actually have an interpretation on its own. And its scale is different depending on your data. For example: what does a deviance of 243.8 mean?

So, we really only use it to compare to the deviances of other models.

Example: Model 2 is a reduced version of Model 1:

$$D_2 - D_1 \sim \chi^2_{p_1 - p_2}$$

An analysis of this can tell you if Model 2 is significantly different from Model 1.



Generalized Pearson's χ^2

In addition to the deviance, there is another measure of model fit, the *generalized Pearson's χ^2* :

$$X^2 = \sum \frac{(y_i - \hat{y}_i)^2}{V(\hat{y}_i)} \sim \chi_p^2$$

Where $V(\mu)$ is the “variance function” from the Aitken EF form. Here it is for a few common GLMs:

	Normal	Poisson	Bernoulli	Gamma
$V(\mu)$	1	μ	$\mu(1 - \mu)$	μ^2



GOTO: Code-a-long: Part I

<https://github.com/TimothyKBook/ga-glm-lecture>

The Exciting Stuff (GLMs for Classification!)

0/1 Binomial Logistic Regression

- You might already be familiar with this.
- Your predictions are on a *probability scale*.

0.32	0.84	0.54	0.12	0.37	0.66	0.03
------	------	------	------	------	------	------

- Can use a cutoff (usually $\frac{1}{2}$), and threshold predictions to be 0/1.


0	1	1	0	0	1	0
---	---	---	---	---	---	---



Example: Fake data from earlier

What is the benefit of using this method?

- It is usually better, or just as good as some of the “fancier” stuff like discriminant analysis and SVMs.
- Computationally cheap.
- Interpretable (can't really explain an SVM, neural net, or RF to “normal people”).



Multiclass Classification

- As I mentioned earlier, the *multinomial distribution* is legal tender for GLMs.
- Recall: The multinomial distribution works just like the binomial distribution, but there are multiple outcomes, rather than just 0/1. It's usually written:

$$Y_i \sim \text{Mult}(n, \mathbf{p}_i)$$

Where \mathbf{p}_i is a vector of probabilities for each class.



Multiclass Classification

So, for a GLM with K possible outcomes, we need to rethink our notation a bit.

$$Y_i \sim \text{ind Mult}(1, \mathbf{p}_i)$$

$$\log \frac{p_1}{p_K} = \beta_{01} + \beta_{11}x_1 + \cdots + \beta_{p1}x_p$$

$$\log \frac{p_2}{p_K} = \beta_{02} + \beta_{12}x_1 + \cdots + \beta_{p2}x_p$$

...

$$\log \frac{p_{K-1}}{p_K} = \beta_{0K-1} + \beta_{1K-1}x_1 + \cdots + \beta_{pK-1}x_p$$

*Notice we now have
 $(p + 1) \times (K - 1)$
parameters to fit!*



Multiclass Classification

For prediction, we solve a system of equations, which your computer can do for you. So, for the first time, *each observation \mathbf{x}_i results in multiple predictions:* $(\hat{p}_1, \dots, \hat{p}_{K-1})$.

Each \hat{p}_j represents the predicted probability that a given \mathbf{x}_i falls into class j .

So, you can *classify* \mathbf{x}_i into which ever has the maximum probability!



Multiclass Classification

Again, in the real world, this actually usually works better than the fancy stuff!

Want proof?

Ghouls, Goblins, and Ghost...

7 months ago · Top 29%

217th

of 764


I used a multinomial logistic regression for my Kaggle submission. I spent about 20 minutes on this. And I even forgot to submit my best one!

I told you I was prolific and world famous.



Ordinal Regression

- Suppose you have a multi-class response again, but this time your data are *ordinal*. That is, they aren't numbers, but they definitely have an order.
 - Example: Small, Medium, Large, X-Large T-shirt sizes.



- How can we tweak multinomial regression to handle this?

Ordinal Regression

- Let's define γ_j to be the *cumulative probability* of being in category j . That is,

$$\gamma_j = \sum_{k=1}^j p_k = P(\text{We are in category} \leq j)$$



The Cumulative Logit Model

Formally, we can define the *cumulative logit model* as:

$$Y_i \sim \text{ind Mult}(1, \mathbf{p})$$

$$\gamma_j = p_1 + \cdots + p_j$$

$$\text{logit } \gamma_1 = \beta_{01} + \beta_{11}x_i + \cdots + \beta_{p1}x_i$$

$$\text{logit } \gamma_2 = \beta_{02} + \beta_{12}x_i + \cdots + \beta_{p2}x_i$$

...

$$\text{logit } \gamma_{K-1} = \beta_{0K-1} + \beta_{1K-1}x_i + \cdots + \beta_{pK-1}x_i$$



And *prediction* can be done recursively:

$$p_1 = \gamma_1$$

$$p_j = \gamma_j - \gamma_{j-1}$$

$$p_K = 1$$

GOTO: Code-a-long Part II

<https://github.com/TimothyKBook/ga-glm-lecture>


GLM Pot Pourri

Regularization

You might remember from last week the concept of *regularization* (or *penalized regression*).

Cliffsnotes: We want to optimize the likelihood, *but not that optimally*. For the *LASSO*, we minimize the following function:

$$\text{stuff} + \lambda \|\beta\|_1$$



Well – you guessed it. You can regularize GLMs.

If you attempted to fit a logistic regression using the `sklearn.LogisticRegression()` method, it actually tried to regularize by default, which is *annoying*.

Otherwise, you can use `sm.GLM().fit_regularized()`.

Generalized Generalized Linear Models

Yes you read that right. What if you think your response variables are actually correlated? You can fit a *Generalized Estimating Equation (GEE)*.

GEEs can be fit in `statsmodels.gee()`.

WARNING: These take a *long* time to fit. Like, more than 10 minutes even for a small data set.



Overdispersion

- Remember that all of these models also assumed *constant variance* among residuals. But...
- If $X \sim \text{Poi}(\lambda)$, then $E[X] = \lambda$ and $\text{Var}[X] = \lambda$
- However, maximum likelihood methods often predict model variance to be *larger* than the response variable's mean.
- This effect can carry over into some GLMs. It's called *overdispersion*, and it's surprisingly common.
- An easy way to detect overdispersion is if the *residual degrees of freedom > deviance*.



- **How can we fix this? There are a few methods:**
 - Some people propose using a different link function. This isn't really a fix, though. But it works sometimes.
 - You can assume a more general variance structure. This can be done with *quasi-likelihood* methods. These are included in `statsmodels.GLM()`. (Similar to GEEs, these take a long time to fit).
 - You can use *Bayesian methods*, if you really want...



Thank you guys for having me! Any Questions?

